

Visualization of Data Movements and Accesses

Til Mohr

Sum of Matrix Elements

Listing 1: Matrix Summation

```
1 let matrix = Matrix::random(2, 2);
2 let mut sum = 0;
3 for column in 0..2 {
4     for row in 0..2 {
5         sum += matrix.get(row, column);
6     }
7 }
8 sum
```

Matrix:

	0	1
0	1	2
1	3	4

Sum of Matrix Elements

Listing 1: Matrix Summation

```
1 let matrix = Matrix::random(2, 2);
2 let mut sum = 0;
3 for column in 0..2 {
4     for row in 0..2 {
5         sum += matrix.get(row, column);
6     }
7 }
8 sum
```

Matrix:

	0	1
0	1	2
1	3	4

Matrix in Memory:

1	2	3	4
---	---	---	---

Sum of Matrix Elements

Listing 1: Matrix Summation

```
1 let matrix = Matrix::random(2, 2);
2 let mut sum = 0;
3 for column in 0..2 {
4     for row in 0..2 {
5         sum += matrix.get(row, column);
6     }
7 }
8 sum
```

Matrix:

	0	1
0	1	2
1	3	4

Current Item:

0	0
---	---

Matrix in Memory:

1	2	3	4
---	---	---	---

Cache:

--	--

Number of cache misses: 0

Sum of Matrix Elements

Listing 1: Matrix Summation

```
1 let matrix = Matrix::random(2, 2);
2 let mut sum = 0;
3 for column in 0..2 {
4     for row in 0..2 {
5         sum += matrix.get(row, column);
6     }
7 }
8 sum
```

Matrix:

	0	1
0	1	2
1	3	4

Current Item:

0	0
---	---

Matrix in Memory:

1	2	3	4
---	---	---	---

Cache:

1	2
---	---

Number of cache misses: 1

Sum of Matrix Elements

Listing 1: Matrix Summation

```
1 let matrix = Matrix::random(2, 2);
2 let mut sum = 0;
3 for column in 0..2 {
4     for row in 0..2 {
5         sum += matrix.get(row, column);
6     }
7 }
8 sum
```

Matrix:

	0	1
0	1	2
1	3	4

Current Item:

1	0
---	---

Matrix in Memory:

1	2	3	4
---	---	---	---

Cache:

1	2
---	---

Number of cache misses: 1

Sum of Matrix Elements

Listing 1: Matrix Summation

```
1 let matrix = Matrix::random(2, 2);
2 let mut sum = 0;
3 for column in 0..2 {
4     for row in 0..2 {
5         sum += matrix.get(row, column);
6     }
7 }
8 sum
```

Matrix:

	0	1
0	1	2
1	3	4

Current Item:

1	0
---	---

Matrix in Memory:

1	2	3	4
---	---	---	---

Cache:

3	4
---	---

Number of cache misses: 2

Sum of Matrix Elements

Listing 1: Matrix Summation

```
1 let matrix = Matrix::random(2, 2);
2 let mut sum = 0;
3 for column in 0..2 {
4     for row in 0..2 {
5         sum += matrix.get(row, column);
6     }
7 }
8 sum
```

Matrix:

	0	1
0	1	2
1	3	4

Current Item:

0	1
---	---

Matrix in Memory:

1	2	3	4
---	---	---	---

Cache:

3	4
---	---

Number of cache misses: 2

Sum of Matrix Elements

Listing 1: Matrix Summation

```
1 let matrix = Matrix::random(2, 2);
2 let mut sum = 0;
3 for column in 0..2 {
4     for row in 0..2 {
5         sum += matrix.get(row, column);
6     }
7 }
8 sum
```

Matrix:

	0	1
0	1	2
1	3	4

Current Item:

0	1
---	---

Matrix in Memory:

1	2	3	4
---	---	---	---

Cache:

1	2
---	---

Number of cache misses: 3

Sum of Matrix Elements

Listing 1: Matrix Summation

```
1 let matrix = Matrix::random(2, 2);
2 let mut sum = 0;
3 for column in 0..2 {
4     for row in 0..2 {
5         sum += matrix.get(row, column);
6     }
7 }
8 sum
```

Matrix:

	0	1
0	1	2
1	3	4

Current Item:

1	1
---	---

Matrix in Memory:

1	2	3	4
---	---	---	---

Cache:

1	2
---	---

Number of cache misses: 3

Sum of Matrix Elements

Listing 1: Matrix Summation

```
1 let matrix = Matrix::random(2, 2);
2 let mut sum = 0;
3 for column in 0..2 {
4     for row in 0..2 {
5         sum += matrix.get(row, column);
6     }
7 }
8 sum
```

Matrix:

	0	1
0	1	2
1	3	4

Current Item:

1	1
---	---

Matrix in Memory:

1	2	3	4
---	---	---	---

Cache:

3	4
---	---

Number of cache misses: 4

Sum of Matrix Elements - Reordered Loops

Listing 2: Matrix Summation

```
1 let matrix = Matrix::random(2, 2);
2 let mut sum = 0;
3 for row in 0..2 {
4     for column in 0..2 {
5         sum += matrix.get(row, column);
6     }
7 }
8 sum
```

Matrix:

	0	1
0	1	2
1	3	4

Current Item:

0	0
---	---

Matrix in Memory:

1	2	3	4
---	---	---	---

Cache:

--	--

Number of cache misses: 0

Sum of Matrix Elements - Reordered Loops

Listing 2: Matrix Summation

```
1 let matrix = Matrix::random(2, 2);
2 let mut sum = 0;
3 for row in 0..2 {
4     for column in 0..2 {
5         sum += matrix.get(row, column);
6     }
7 }
8 sum
```

Matrix:

	0	1
0	1	2
1	3	4

Current Item:

0	0
---	---

Matrix in Memory:

1	2	3	4
---	---	---	---

Cache:

1	2
---	---

Number of cache misses: 1

Sum of Matrix Elements - Reordered Loops

Listing 2: Matrix Summation

```
1 let matrix = Matrix::random(2, 2);
2 let mut sum = 0;
3 for row in 0..2 {
4     for column in 0..2 {
5         sum += matrix.get(row, column);
6     }
7 }
8 sum
```

Matrix:

	0	1
0	1	2
1	3	4

Matrix in Memory:

1	2	3	4
---	---	---	---

Current Item:

0	1
---	---

Cache:

1	2
---	---

Number of cache misses: 1

Sum of Matrix Elements - Reordered Loops

Listing 2: Matrix Summation

```
1 let matrix = Matrix::random(2, 2);
2 let mut sum = 0;
3 for row in 0..2 {
4     for column in 0..2 {
5         sum += matrix.get(row, column);
6     }
7 }
8 sum
```

Matrix:

	0	1
0	1	2
1	3	4

Current Item:

1	0
---	---

Matrix in Memory:

1	2	3	4
---	---	---	---

Cache:

1	2
---	---

Number of cache misses: 1

Sum of Matrix Elements - Reordered Loops

Listing 2: Matrix Summation

```
1 let matrix = Matrix::random(2, 2);
2 let mut sum = 0;
3 for row in 0..2 {
4     for column in 0..2 {
5         sum += matrix.get(row, column);
6     }
7 }
8 sum
```

Matrix:

	0	1
0	1	2
1	3	4

Current Item:

1	0
---	---

Matrix in Memory:

1	2	3	4
---	---	---	---

Cache:

3	4
---	---

Number of cache misses: 2

Sum of Matrix Elements - Reordered Loops

Listing 2: Matrix Summation

```
1 let matrix = Matrix::random(2, 2);
2 let mut sum = 0;
3 for row in 0..2 {
4     for column in 0..2 {
5         sum += matrix.get(row, column);
6     }
7 }
8 sum
```

Matrix:

	0	1
0	1	2
1	3	4

Current Item:

1	1
---	---

Matrix in Memory:

1	2	3	4
---	---	---	---

Cache:

3	4
---	---

Number of cache misses: 2

Outline

- Memory-Related Performance Problems
 - Data Locality
 - Processor-Memory Performance Gap
- Overview of the Optimization Workflow
- Data Gathering Approaches
- Visualization Techniques
- Specific Optimization Tool
- Conclusion

Memory-Related Performance Problems | Data Locality

$$t_{avg} = p \cdot t_c + (1 - p) \cdot t_m \quad (1)$$

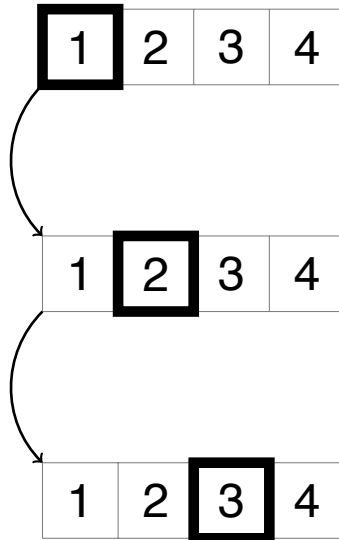
t_{avg} :	average access time
p :	cache hit percentage
t_c :	cache access time
t_m :	memory access time

$$t_c \ll t_m \quad (2)$$

Memory-Related Performance Problems I Data Locality

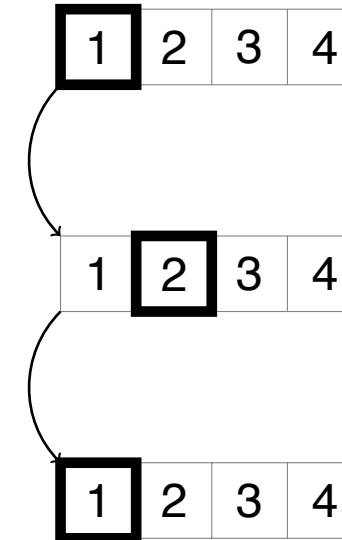
Spatial Locality

- Data that is referenced spatially close together is likely to be accessed in the near future

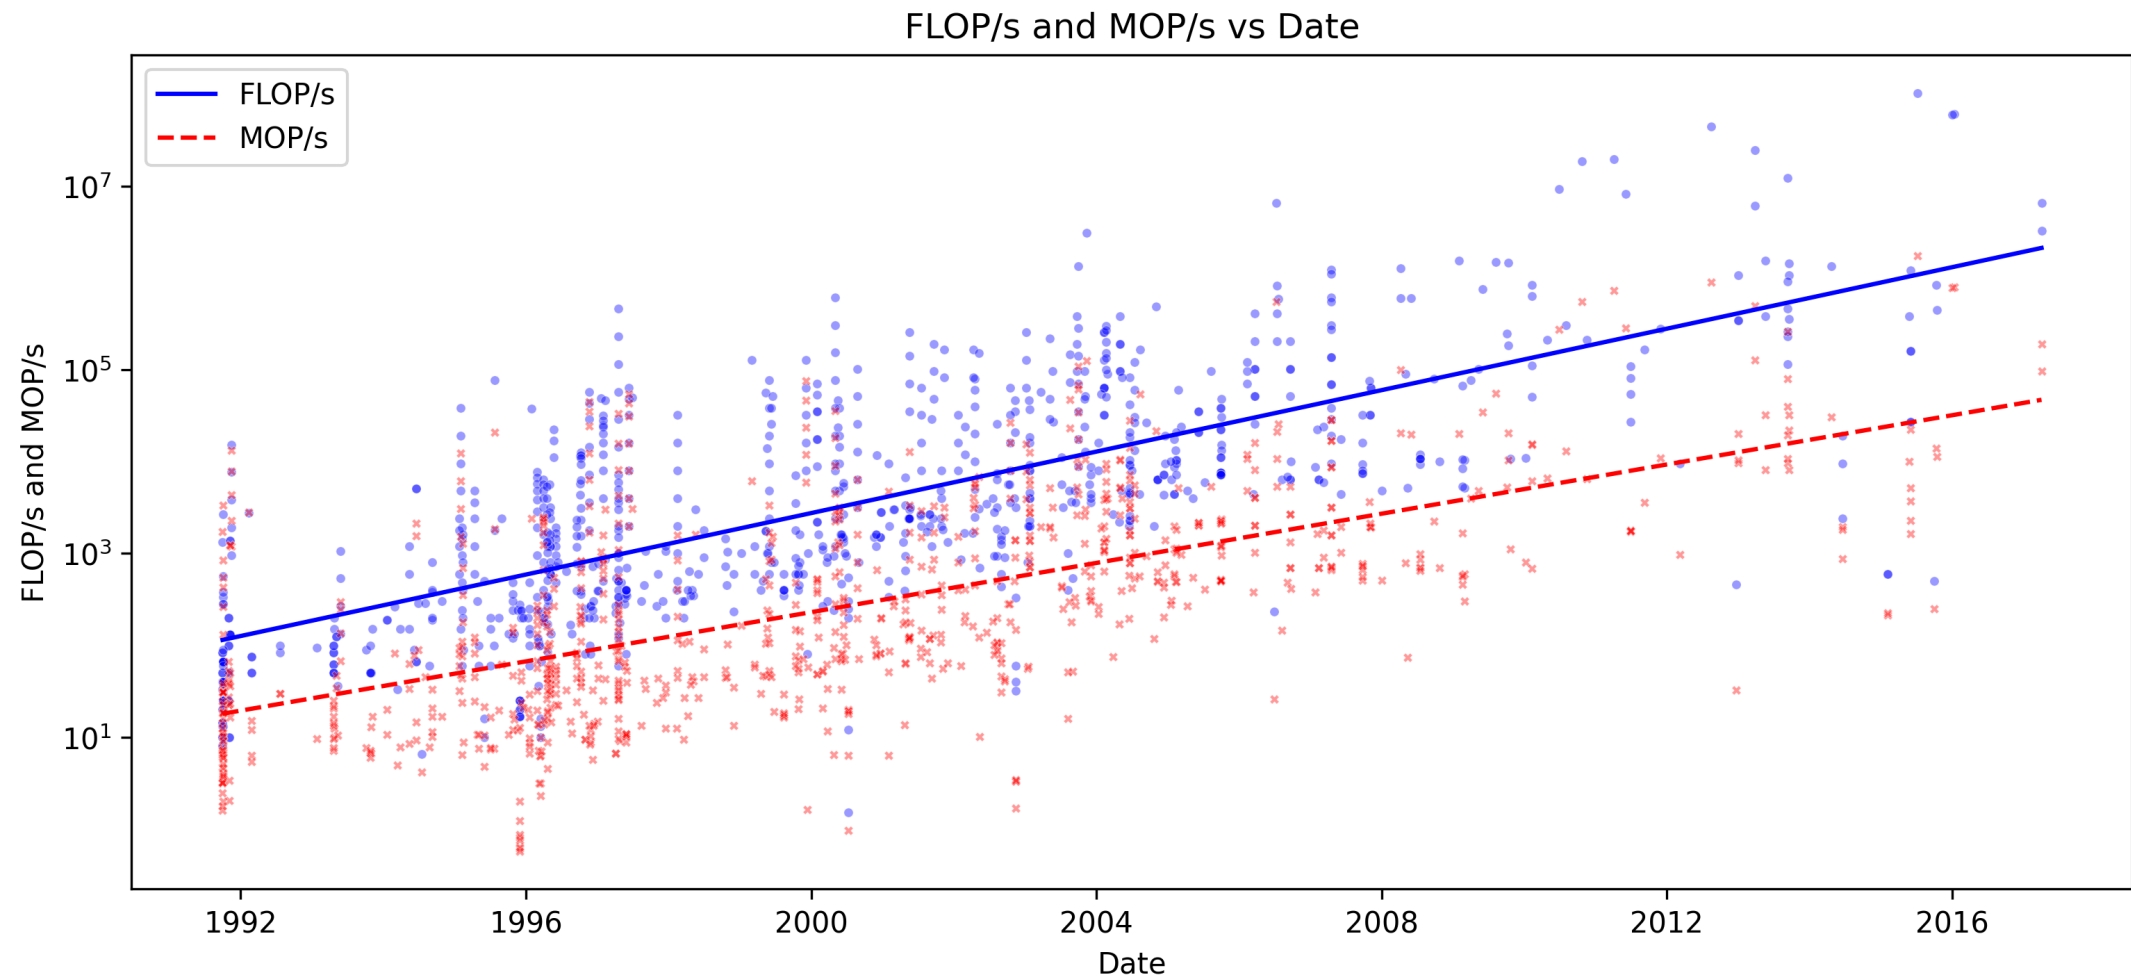


Temporal Locality

- Data that is referenced in the near past is likely to be accessed in the near future

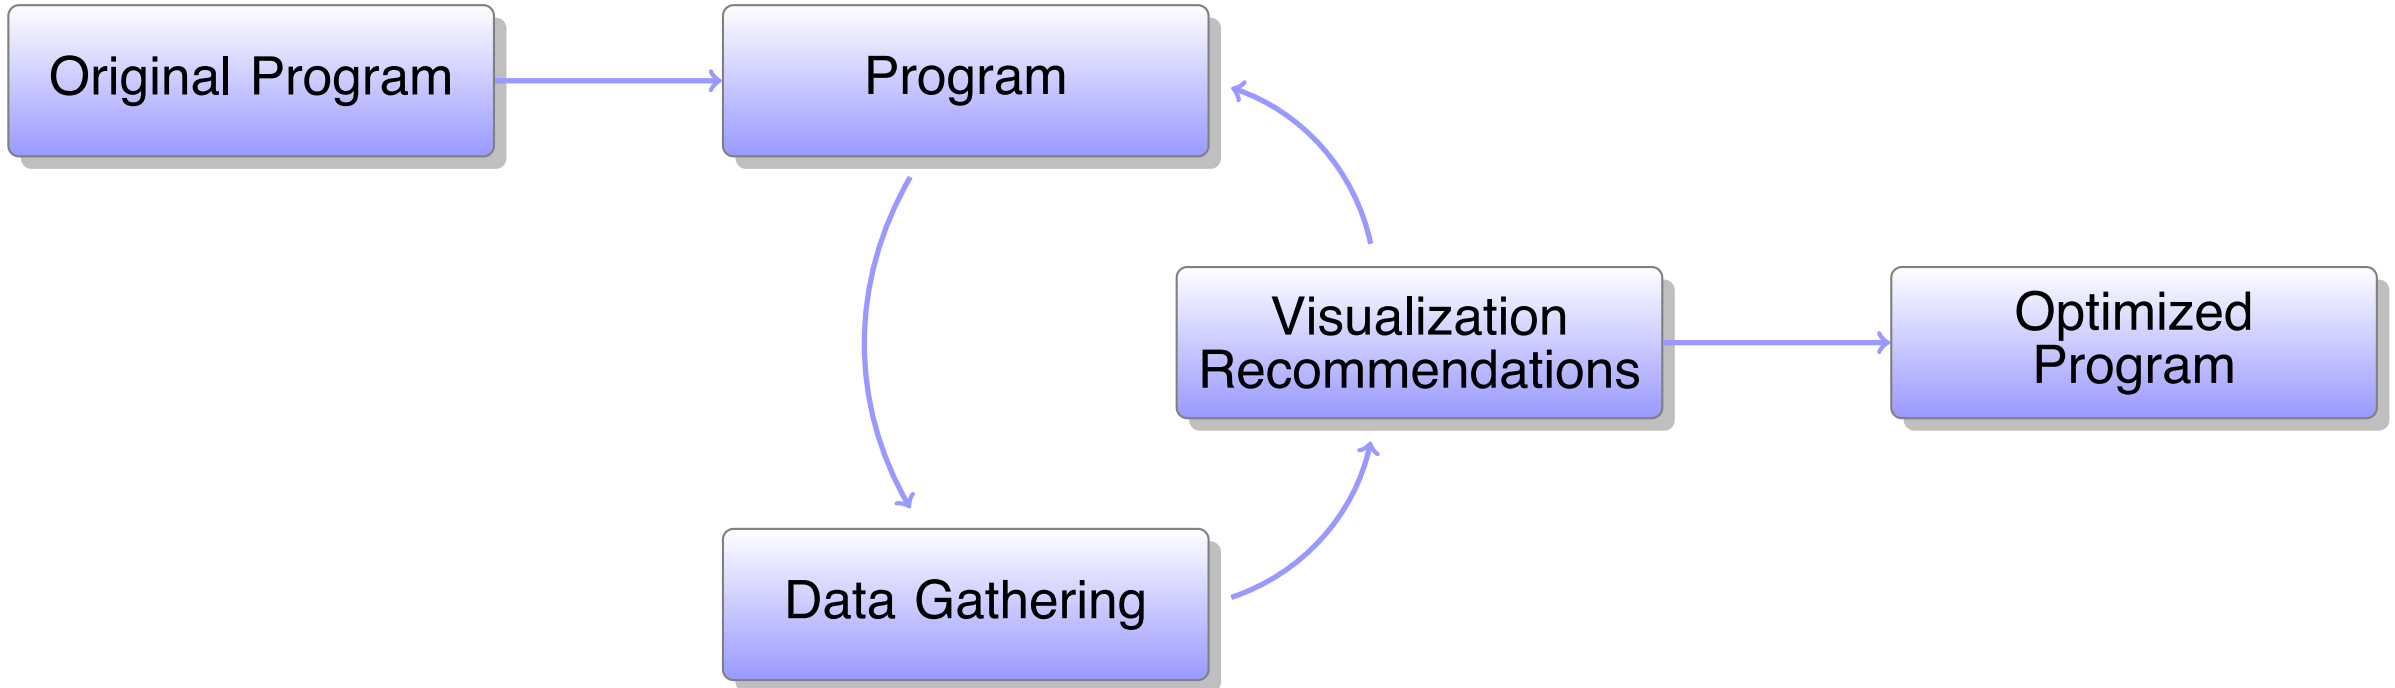


Memory-Related Performance Problems | Processor-Memory Performance Gap



Overview of the Optimization Workflow

Visualization-Guided Optimization



Goal: Acquire Memory-Related Data for Visualization

- Data Accesses
 - Memory locations / variables
 - Frequencies
- Data access patterns
 - Nested loops
- Cache performance
 - Hit/miss rates
 - Utilization
 - Amount of data transfer in between different cache levels and main memory

Run Program and Capture Memory-Related Information

- Hardware counters
 - Counts cache hits/misses
- Tracing / profiling
- Store source code references alongside memory-related information


- 😊 Very accurate
 - Real program data
 - Actual physical hardware

- 😞 Can be very slow
- 😞 Possible large overhead for very granular data
- 😞 Cannot easily analyze just parts of the program

Analyze the Programs Source Code for Data Accesses

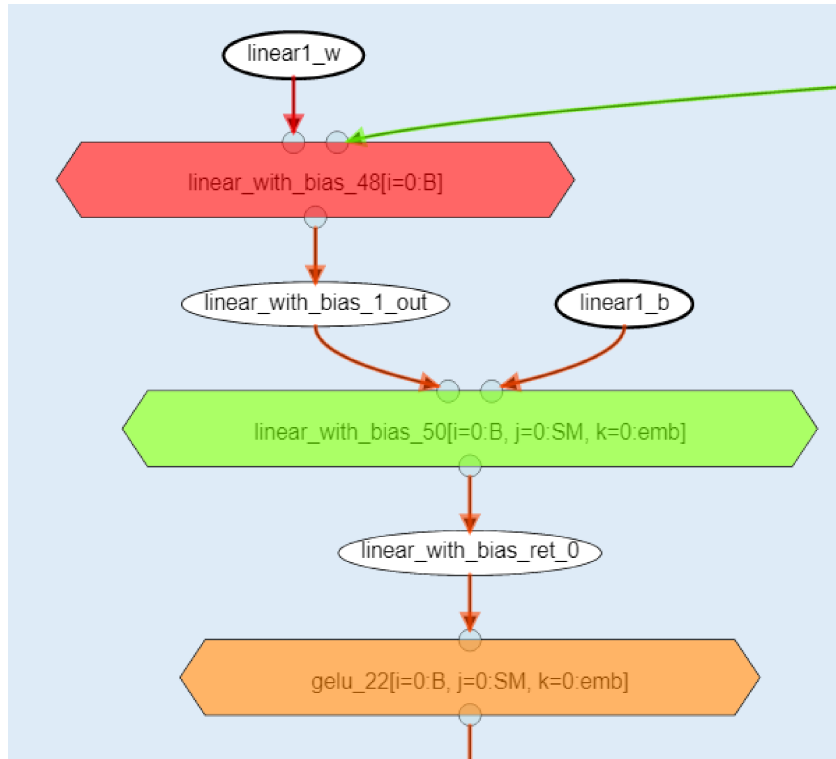
- Extract any data access information purely from the source code
- Compile the program into a Data-Flow Oriented IR
- Statistics gathered by analyzing the IR
 - Algorithmic intensity
 - Volume of data circulating in the program

```
void square(double &v) {  
    v = v * v;  
}  
  
// ...  
  
square(A[i]);
```



Source: Alexandru Calotoiu et al. "Lifting C semantics for dataflow optimization". In: *Proceedings of the 36th ACM International Conference on Supercomputing*. 2022, pp. 1-13.

Analyze the Programs Source Code for Data Accesses



Source: Alexandru Calotoiu et al. "Lifting C semantics for dataflow optimization". In: *Proceedings of the 36th ACM International Conference on Supercomputing*. 2022, pp. 1-13.

- 😊 Very fast
- 😊 Provides holistic view of the program and its performance
- 😞 Very abstract analysis
 - Memory layout of data is not considered
 - Hardware architecture unknown
 - No information about real-world cache performance

Imitate the Programs Memory Accesses on a Simulated Cache Hierarchy

References I

- 📖 Alexandru Calotoiu et al. “Lifting C semantics for dataflow optimization”. In: *Proceedings of the 36th ACM International Conference on Supercomputing*. 2022, pp. 1-13.